

# Class 1 Notes

## Combining visualization and transformation (r4ds - Section I)

For the first 2 classes, we will take the tools in r4ds section 1 - Explore, and walk through scenarios:

Load the tidyverse library, and get the mpg data (*included in the base package*). Put that in a dataframe. Look at the structure of the data, and the first 10 rows:

```
library(tidyverse)
```

```
str(mpg) # this gives you the structure of the dataframe
```

```
Classes 'tbl_df', 'tbl' and 'data.frame':  234 obs. of  11 variables:
 $ manufacturer: chr  "audi" "audi" "audi" "audi" ...
 $ model       : chr  "a4" "a4" "a4" "a4" ...
 $ displ      : num  1.8 1.8 2 2 2.8 2.8 3.1 1.8 1.8 2 ...
 $ year       : int  1999 1999 2008 2008 1999 1999 2008 1999 1999 2008 ...
 $ cyl        : int   4 4 4 4 6 6 6 4 4 4 ...
 $ trans      : chr  "auto(l5)" "manual(m5)" "manual(m6)" "auto(av)" ...
 $ drv        : chr  "f" "f" "f" "f" ...
 $ cty        : int  18 21 20 21 16 18 18 18 16 20 ...
 $ hwy        : int  29 29 31 30 26 26 27 26 25 28 ...
 $ fl         : chr  "p" "p" "p" "p" ...
 $ class      : chr  "compact" "compact" "compact" "compact" ...
```

```
top_n(mpg, 10)
```

```
# A tibble: 62 x 11
  manufacturer model      displ  year   cyl trans drv     cty   hwy fl      class
    <chr>         <chr>    <dbl> <int> <int> <chr> <chr> <int> <int> <chr> <chr>
1 chevrolet    c1500 sub...   5.3  2008     8 auto... r       14    20 r      suv
2 chevrolet    c1500 sub...   5.3  2008     8 auto... r       11    15 e      suv
3 chevrolet    c1500 sub...   5.3  2008     8 auto... r       14    20 r      suv
4 chevrolet    c1500 sub...   5.7  1999     8 auto... r       13    17 r      suv
5 chevrolet    c1500 sub...   6    2008     8 auto... r       12    17 r      suv
6 chevrolet    k1500 tah...   5.3  2008     8 auto... 4       14    19 r      suv
7 chevrolet    k1500 tah...   5.3  2008     8 auto... 4       11    14 e      suv
8 chevrolet    k1500 tah...   5.7  1999     8 auto... 4       11    15 r      suv
9 chevrolet    k1500 tah...   6.5  1999     8 auto... 4       14    17 d      suv
10 dodge       durango 4...   3.9  1999     6 auto... 4       13    17 r      suv
# ... with 52 more rows
```

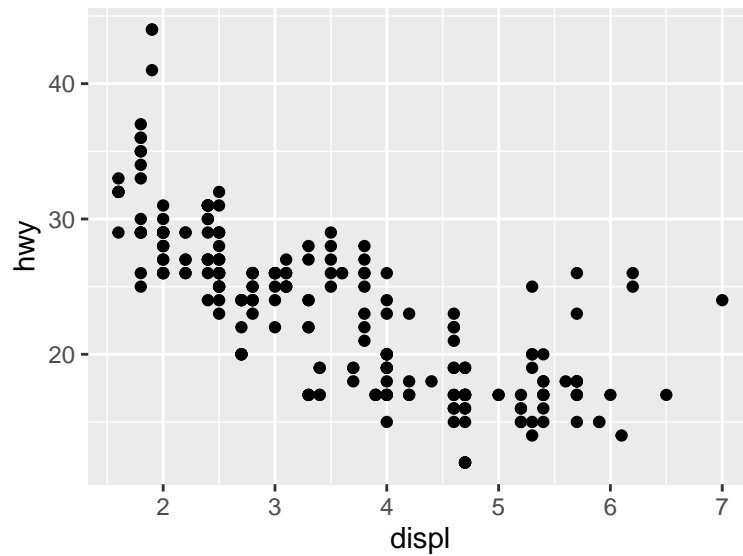
What do the datatypes in str tell you. What's an int? a chr? a num?

Now that we have some data, let's work with visualization in ggplot:

(Note that I use a different syntax from Hadley - I store the plots in "objects" - this time I named it "p", but you can name it anything. Then I can add whatever I like later without having to recreate the whole plot). Also, I use a less verbose syntax. Either way is fine - up to you.

Note: with the aesthtics (*aes*), x comes first.

```
p = ggplot(mpg, aes(displ,hwy))
p = p + geom_point()
p
```

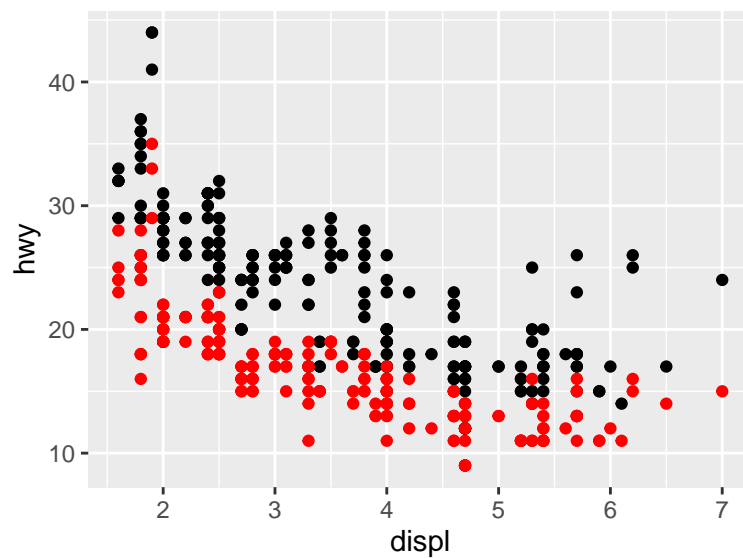


So, what does this tell us? Does it make sense? If it doesn't, go back and figure out. It's not about drawing pictures, it's about understanding the data. (*Bots can draw pictures - do you want to compete with bots?*)

Now that we have highway mileage, let's add city and compare (*notice how we are just adding a **layer** - one of the beauties of ggplot - it really is well designed*).

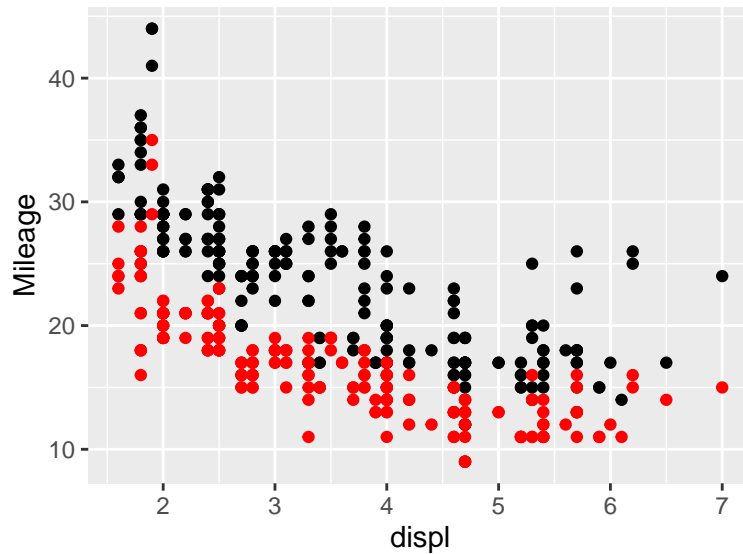
Does this all make sense? Always ask! So, we'll add a new layer and color it red:

```
p = p + geom_point(aes(displ, cty), color = "red")
p
```



There's something wrong with this plot, though - can you see it? it's not hwy mileage anymore, let's correct it:

```
p <- p + ylab("Mileage")
p
```

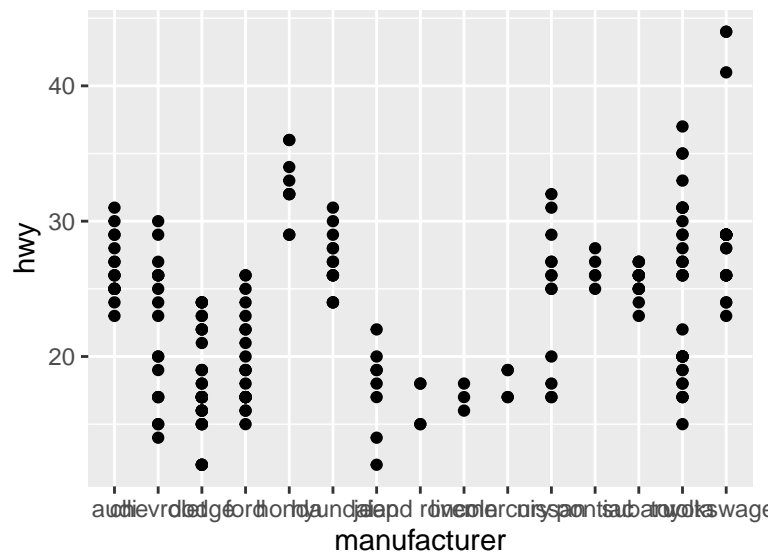


Notice again how we didn't have to recreate the whole plot - we just added a layer.

What if we wanted to see mileage by manufacturer? How would we do it? **Think**

Let's put manufacturer on the X axis instead of displ:

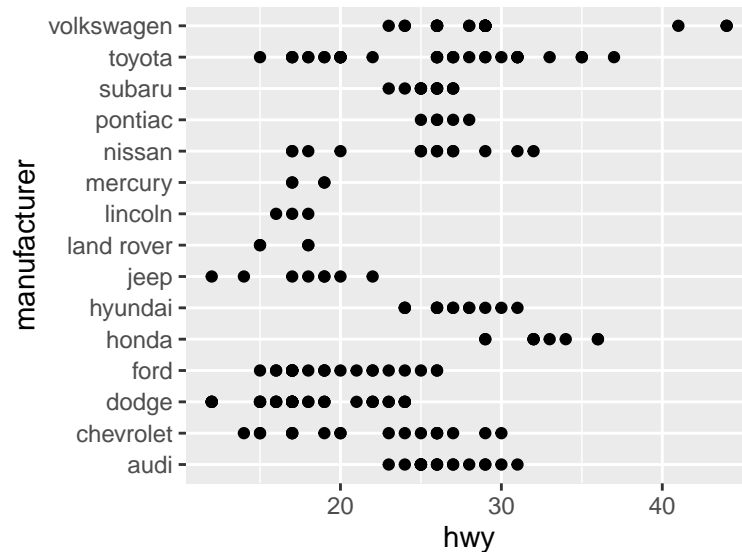
```
p = ggplot(mpg, aes(manufacturer,hwy))
p = p + geom_point()
p
```



What does each of the points represent? You should know. Also, notice that we DID need to create a new plot here - we're no longer building.

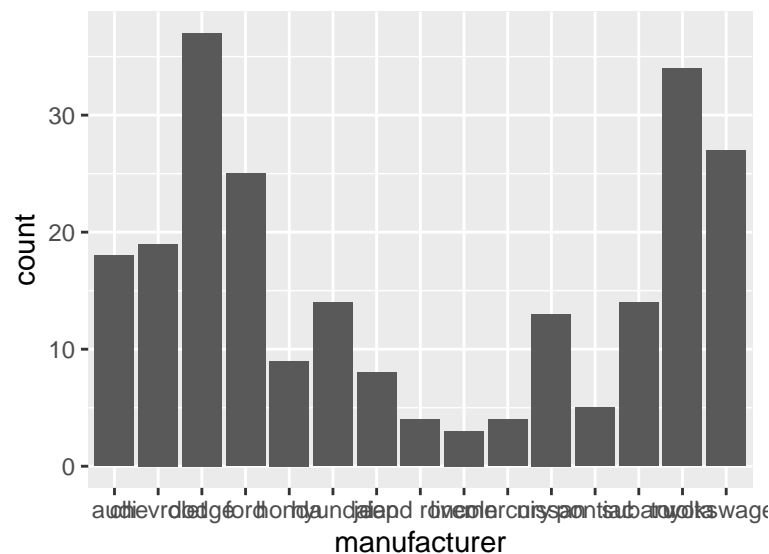
This is hard to read. We can fix that by flipping the axis (*add a layer*):

```
p = p + coord_flip()
p
```



This might look better as a bar *(we'll create a new plot)*

```
p = ggplot(mpg, aes(manufacturer)) + geom_bar()
p
```



Wait! what is this? Does this make sense? **THINK**. `geom_bar` defaults to count. Is that what we want?

We want to know hwy mileage by maufacturer, but the sample had a different number of cars by manufacturer *(look at it - know your data)* We can verify that with the function `count`.

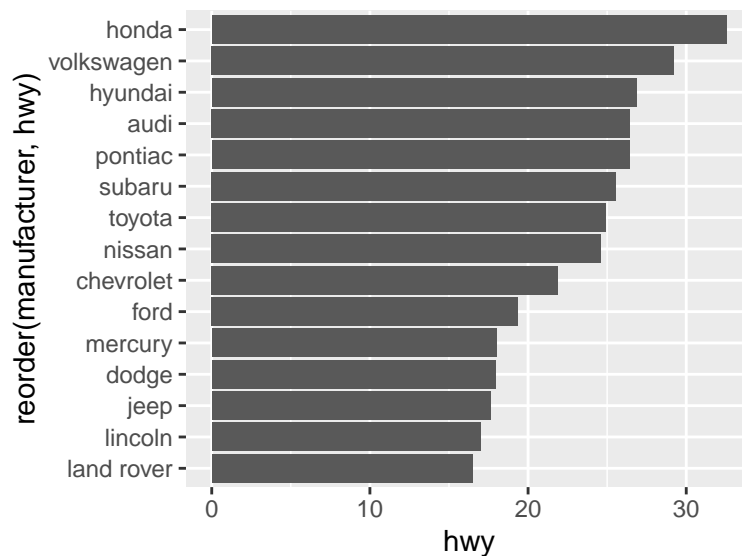
```
count(mpg, manufacturer)
```

```
# A tibble: 15 x 2
  manufacturer     n
  <chr>         <int>
1 audi           18
2 chevrolet       19
3 dodge           36
4 ford           25
5 honda           9
6 hyundai        14
7 jeep            8
8 land rover      4
9 lincoln         3
10 mercury         4
11 nissan          13
12 pontiac         5
13 subaru         14
14 toyota         34
15 volkswagen     27
```

1	audi	18
2	chevrolet	19
3	dodge	37
4	ford	25
5	honda	9
6	hyundai	14
7	jeep	8
8	land rover	4
9	lincoln	3
10	mercury	4
11	nissan	13
12	pontiac	5
13	subaru	14
14	toyota	34
15	volkswagen	27

So, is that what we want to know? **NO!** We want to see mileage, right? Here's one way:

```
p = ggplot(mpg) +
  geom_bar(aes(reorder(manufacturer, hwy), hwy), stat = "summary", fun.y = "mean") +
  coord_flip()
p
```



A little explanation:

**reorder(manufacturer, hwy)** reorders the x axis by hwy (*that way, we can see the ranking*)

**stat = "summary", fun.y = "mean"** tells ggplot that we want a summary function and that function is the mean. (*geom\_bar and geom\_histogram will normally default to count. This command will override that - but you can also do your own math here*)

That said, This is not usually the way you would build this in practice. That's because you are usually working with a particular stat, and you're going to continue your analysis beyond the exploratory visualization. So, normally, you would create the stat, and **store** it, so you can use it later - something like:

```
mpg %>% group_by(manufacturer) %>% summarise(AvgHwy = mean(hwy))
```

```
# A tibble: 15 x 2
  manufacturer AvgHwy
```

	<chr>	<dbl>
1	audi	26.4
2	chevrolet	21.9
3	dodge	17.9
4	ford	19.4
5	honda	32.6
6	hyundai	26.9
7	jeep	17.6
8	land rover	16.5
9	lincoln	17
10	mercury	18
11	nissan	24.6
12	pontiac	26.4
13	subaru	25.6
14	toyota	24.9
15	volkswagen	29.2

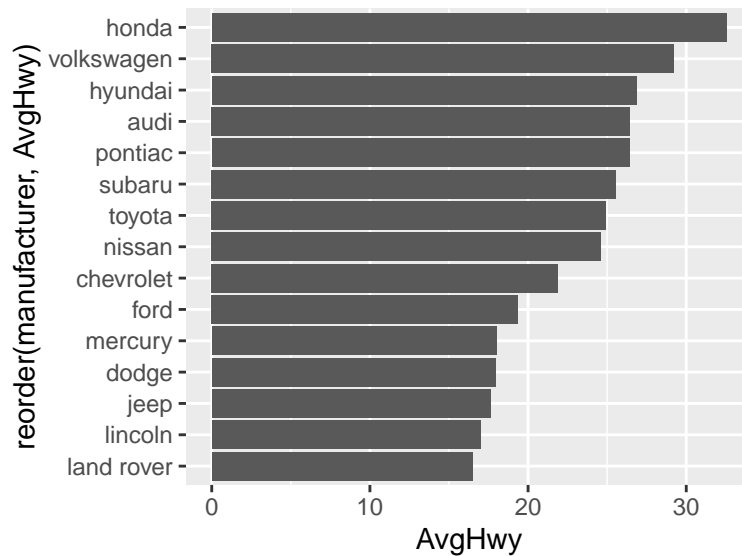
And then we can plot that. BUT, you have to remember to save (*persist*) the analysis to an object (*it hasn't been saved - you just asked R to do a group\_by AND THEN do a summarise, but you did not save it*). To save it (*and use it later*), you can:

```
dfAvgHwy = mpg %>% group_by(manufacturer) %>% summarise(AvgHwy = mean(hwy))
head(dfAvgHwy)
```

```
# A tibble: 6 x 2
  manufacturer AvgHwy
  <chr>         <dbl>
1 audi         26.4
2 chevrolet    21.9
3 dodge        17.9
4 ford         19.4
5 honda        32.6
6 hyundai      26.9
```

And then plot it:

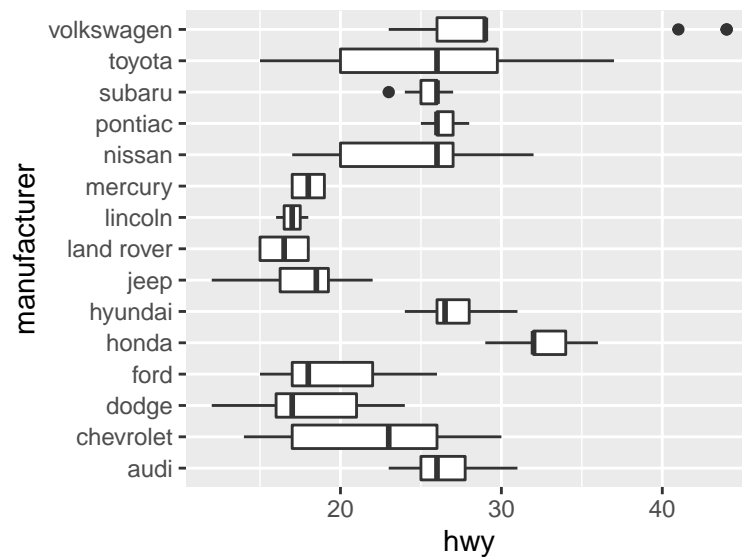
```
p = ggplot(dfAvgHwy, aes(reorder(manufacturer, AvgHwy), AvgHwy)) +
  geom_bar(stat = "identity") +
  coord_flip()
p
```



Notice how we have changed datasource and y axis (*which becomes the x axis after we do a coordinate flip*). Please understand this process. This is a very basic and expected skill.

There are better ways to look at all this. Boxplots are great:

```
p = ggplot(data = mpg, aes(x = manufacturer, y = hwy)) +
  geom_boxplot() +
  coord_flip()
p
```



The boxplot gives us a visual of the distribution of data by manufacturer. It also shows the mean, AND the first and third quartiles.

Look at volkswagen. That should get your attention - you should notice that the mean and third quartile are really close, with 2 “outliers” beyond 40 hwy. Let’s take a summary look at the data:

```
summary(filter(mpg, manufacturer == "volkswagen")$hwy)
```

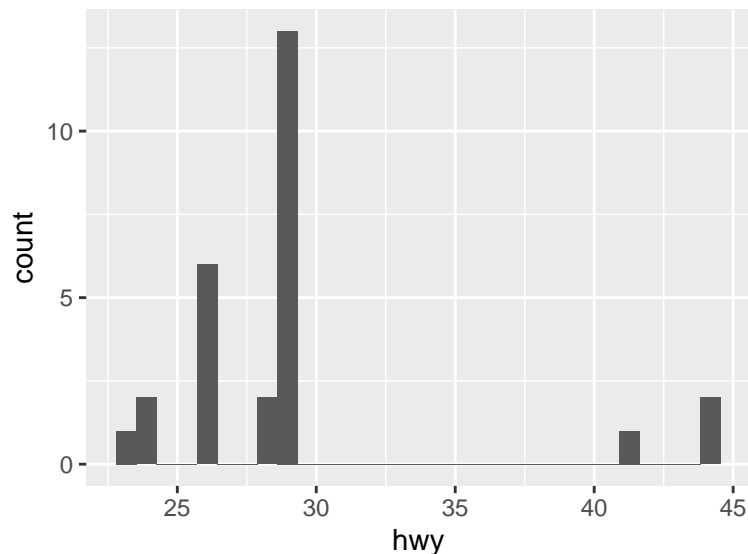
Min. 1st Qu. Median Mean 3rd Qu. Max.

23.00 26.00 29.00 29.22 29.00 44.00

(summary can be applied to a range of objects - we'll use it extensively)

So far, its hard to tell why the mean is right next to the 3rd quartile. Let's look at a histogram:

```
p = ggplot(filter(mpg, manufacturer == "volkswagen"), aes(hwy)) + geom_histogram()
p
```



What does a quantile, or quartile mean? It's how the data are ordered. What is a mean? Could it be that there's a group of observations that are "pulling" the mean up - skewing it? Let's take a closer look:

```
filter(mpg, manufacturer == "volkswagen") %>% arrange(desc(hwy))
```

# A tibble: 27 x 11

	manufacturer	model	displ	year	cyl	trans	drv	cty	hwy	fl	class
	<chr>	<chr>	<dbl>	<int>	<int>	<chr>	<chr>	<int>	<int>	<chr>	<chr>
1	volkswagen	jetta	1.9	1999	4	manual...	f	33	44	d	compa...
2	volkswagen	new be...	1.9	1999	4	manual...	f	35	44	d	subco...
3	volkswagen	new be...	1.9	1999	4	auto(l...	f	29	41	d	subco...
4	volkswagen	gti	2	1999	4	manual...	f	21	29	r	compa...
5	volkswagen	gti	2	2008	4	manual...	f	21	29	p	compa...
6	volkswagen	gti	2	2008	4	auto(s...	f	22	29	p	compa...
7	volkswagen	jetta	2	1999	4	manual...	f	21	29	r	compa...
8	volkswagen	jetta	2	2008	4	auto(s...	f	22	29	p	compa...
9	volkswagen	jetta	2	2008	4	manual...	f	21	29	p	compa...
10	volkswagen	jetta	2.5	2008	5	auto(s...	f	21	29	r	compa...

# ... with 17 more rows

Comparing jettas and new beetles

```
filter(mpg, manufacturer == "volkswagen", model %in% c("jetta", "new beetle"))
```

# A tibble: 15 x 11

	manufacturer	model	displ	year	cyl	trans	drv	cty	hwy	fl	class
	<chr>	<chr>	<dbl>	<int>	<int>	<chr>	<chr>	<int>	<int>	<chr>	<chr>
1	volkswagen	jetta	1.9	1999	4	manual...	f	33	44	d	compa...
2	volkswagen	jetta	2	1999	4	manual...	f	21	29	r	compa...
3	volkswagen	jetta	2	1999	4	auto(l...	f	19	26	r	compa...



4	volkswagen	jetta	2	2008	4	auto(s... f	22	29	p	compa...
5	volkswagen	jetta	2	2008	4	manual... f	21	29	p	compa...
6	volkswagen	jetta	2.5	2008	5	auto(s... f	21	29	r	compa...
7	volkswagen	jetta	2.5	2008	5	manual... f	21	29	r	compa...
8	volkswagen	jetta	2.8	1999	6	auto(l... f	16	23	r	compa...
9	volkswagen	jetta	2.8	1999	6	manual... f	17	24	r	compa...
10	volkswagen	new be...	1.9	1999	4	manual... f	35	44	d	subco...
11	volkswagen	new be...	1.9	1999	4	auto(l... f	29	41	d	subco...
12	volkswagen	new be...	2	1999	4	manual... f	21	29	r	subco...
13	volkswagen	new be...	2	1999	4	auto(l... f	19	26	r	subco...
14	volkswagen	new be...	2.5	2008	5	manual... f	20	28	r	subco...
15	volkswagen	new be...	2.5	2008	5	auto(s... f	20	29	r	subco...

So, These “outliers” are not really outliers. They’re diesels. And that would explain the wide gap.

```
filter(mpg, manufacturer == "volkswagen", model %in% c("jetta", "new beetle"))
```

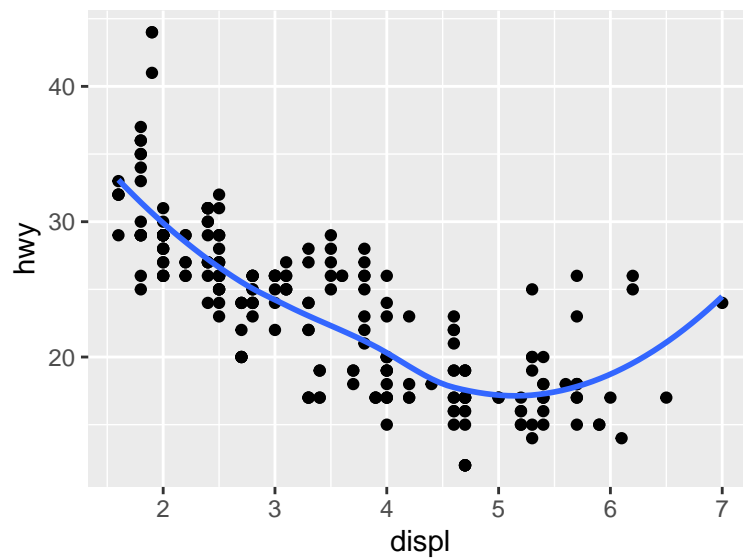
# A tibble: 15 x 11

	manufacturer	model	displ	year	cyl	trans	drv	cty	hwy	fl	class
	<chr>	<chr>	<dbl>	<int>	<int>	<chr>	<chr>	<int>	<int>	<chr>	<chr>
1	volkswagen	jetta	1.9	1999	4	manual... f		33	44	d	compa...
2	volkswagen	jetta	2	1999	4	manual... f		21	29	r	compa...
3	volkswagen	jetta	2	1999	4	auto(l... f		19	26	r	compa...
4	volkswagen	jetta	2	2008	4	auto(s... f		22	29	p	compa...
5	volkswagen	jetta	2	2008	4	manual... f		21	29	p	compa...
6	volkswagen	jetta	2.5	2008	5	auto(s... f		21	29	r	compa...
7	volkswagen	jetta	2.5	2008	5	manual... f		21	29	r	compa...
8	volkswagen	jetta	2.8	1999	6	auto(l... f		16	23	r	compa...
9	volkswagen	jetta	2.8	1999	6	manual... f		17	24	r	compa...
10	volkswagen	new be...	1.9	1999	4	manual... f		35	44	d	subco...
11	volkswagen	new be...	1.9	1999	4	auto(l... f		29	41	d	subco...
12	volkswagen	new be...	2	1999	4	manual... f		21	29	r	subco...
13	volkswagen	new be...	2	1999	4	auto(l... f		19	26	r	subco...
14	volkswagen	new be...	2.5	2008	5	manual... f		20	28	r	subco...
15	volkswagen	new be...	2.5	2008	5	auto(s... f		20	29	r	subco...

We might consider segregating the fuel. Worry about that later.

One last thing, let’s look at a regression line and see if there’s a trend:

```
p = ggplot(mpg, aes(displ,hwy)) +
  geom_point() +
  geom_smooth(se = F)
p
```



We'll get into regression modeling this semester, but for now, you can see there's a trend where hwy mileage decreases with the size of the engine (*displ*). Duh.